

Contact: David Zhang – dzhang [at] cs [dot] toronto [dot] edu

1. The VeB queue works with giant arrays; the time cost of initializing them to all zeros would be prohibitive. Devise a way to avoid initializing large arrays. More specifically, develop a data structure that holds n items according to an index $i \in \{1, \dots, n\}$ and supports the following operations in $O(1)$ time (worst case; i.e. not amortized, not expected) per operation:

Init Initializes the data structure (assuming that the necessary space has been allocated) to empty.

Set(i, x) places item x at index i in the data structure.

Get(i) returns the item at index i , or “empty” if nothing is there.

Your data structure should use $O(n)$ space and should work **regardless** of what garbage values are stored in that space at the beginning of execution. You can assume n fits in one machine word.

Hint: Use extra space to remember which entries of the array have been initialized. But remember: the extra space also starts out with garbage entries!¹

If we desire $O(1)$ worst-case operations, then we cannot use lazy or randomized data structures. We will use arrays to map indices to keys. The big question is: Is the key at index i garbage? This is impossible to answer if using a single array. We can try using two arrays with a *count* variable corresponding to the number of keys that are currently present.

Have one array store all the keys in an ‘active’ subarray from indices $1 \dots \text{count}$. Then, store the position of each key at index i in the second array. When we query the data structure, if index i in the second array returns a position $j \leq \text{count}$, then the first array at position j is a valid key. Unfortunately, we cannot verify that j is not garbage. Therefore, use a third array.

Call the array mapping input indices to the ‘active’ block array, A . Call the array maintaining the ‘active’ block B , and call the array that actually contains the keys C .

- On initialization, allocate space for A, B and C .
- On inserts, assign $C(i) = x, \text{count} = \text{count} + 1, B(\text{count}) = i, A(i) = \text{count}$.
- On queries, first check that $A(i) \leq \text{count}$ to verify that $B(A(i))$ is a valid entry. Further, $B(A(i)) = i$ verifies that the data structure indeed stores the key corresponding to index i . Hence, $C(i)$ is guaranteed to be non-garbage. If $A(i) > \text{count}$ or $B(A(i)) \neq i$, then the data structure does not store anything at index i .

All operations run in $O(1)$ worst-case time and the space requirement is $O(n)$.

¹This question was taken from Karger’s 6.854 Advanced Algorithms Homework:
https://nb.mit.edu/nb_viewer.html?id=4c730237d23143602d1ff49aab3f0d41